
CSCI 4061

Introduction to Operating Systems

Instructor: Abhishek Chandra

Outline

- Socket Overview
- Socket Types
- Socket Operations
- Client-Server Sockets

2

What are Sockets?

- Networking API provided by the OS
 - Enable applications to "hook" onto the network
 - Support one end of a network connection
- Special files
 - Program is given socket file descriptors
 - Read, write data to them
 - Special operations needed

3

Socket Types

- Protocol families
 - E.g.: TCP/IP, UUCP, Appletalk
- Connection-oriented: TCP Sockets
- Connectionless: UDP Sockets

4

Socket Addresses

- What addresses do you need for communicating between two processes on two hosts?
- Each socket is associated with a 5-tuple
 - {protocol, local-addr, local-port, remote-addr, remote-port}

5

Socket Address: struct sockaddr

- Generic type: struct sockaddr
 - Cast to/from the actual structure type
- TCP/IP: struct sockaddr_in
 - Family: AF_INET/PF_INET
 - Port: 16-bit TCP/UDP port number
 - Address: 32-bit IP address structure

```
struct in_addr {
    in_addr_t s_addr;
};
```

6

Address Conversion

- inet_addr: Dotted to binary format
 - E.g.: 192.168.10.3 to its 32-bit equivalent
- inet_ntoa: Reverse function
- gethostbyname: Hostname to IP address
 - Uses DNS

7

Machine Byte Ordering

- Different machines use different byte orders
 - Big-endian: MSB at low addresses
 - E.g.: Sparc, PowerPC
 - Little-endian: MSB at high addresses
 - E.g.: x86
- What if sender is big-endian while receiver is little-endian?

8

Network Byte Order

- Common byte order used for all network data transmission
 - Big-endian
 - Opposite of x86 machine architecture
- Data being sent out on the network must be converted to network byte order and vice versa
- Port numbers and IP addresses should also be converted
 - Why?

9

Network Byte Order Conversion

- Host-to-network byte order conversion
 - htons: 16-bit conversion
 - htonl: 32-bit conversion
- Reverse conversion
 - ntohs, ntohl
- Why don't we have htonc, ntohc?
- Should we use conversion functions if we are programming on a big-endian machine?

10

Socket Operations

- Dependent on
 - Connectionless vs. connection-oriented
 - Client vs. server
- Some generic operations
 - Socket creation
 - Data I/O

11

TCP Sockets: Server Operations

- Create a socket
- Bind a local address/port number
- Wait for connections from clients
- Accept a connection
- Read request, service it, return results
- Close client connection

12

TCP Sockets: Client Operations

- Create a socket
- Connect to a remote server
- Send request
- Receive results
- Close socket

13

Creating a Socket: socket

```
int socket(int family, int type, int protocol);
```

- Returns a file descriptor
 - Identifier for the socket
- Parameters
 - family: Protocol family. E.g.: Internet or TCP/IP (AF_INET/PF_INET)
 - type: Protocol type
 - SOCK_STREAM: Stream (TCP)
 - SOCK_DGRAM: Datagram (UDP)
 - SOCK_RAW: Raw (IP)
 - protocol: Typically 0

14

Binding a Local Address/Port: bind

```
int bind(int sockfd, struct sockaddr *myaddr,  
        socklen_t addrlen);
```

- Parameters:
 - sockfd: socket file descriptor
 - myaddr: address structure containing local address/port
 - addrlen: Length of address structure
- Binds a local address/port based on values specified in the address structure

15

Local Ports

- Well-known ports:
 - Process specifies a non-zero port number
 - Servers typically do this
 - E.g.: Web: 80, ftp: 21, ssh: 22
- Ephemeral ports:
 - Port number specified in address struct is 0
 - Kernel chooses an unused port number from a range
 - Clients typically do this

16

Local Addresses

- Process may specify a local IP address
 - One of the valid network interface addresses
 - Communication would happen through the chosen address
- Process may specify a wildcard IP address
 - INADDR_ANY
 - Kernel will choose a default IP address

17

Listening for Connections: listen

```
int listen(int sockfd, int backlog);
```

- Converts a socket to a “passive” server socket
 - Called only by a TCP server
- Parameters:
 - sockfd: socket file descriptor
 - backlog: Number of pending client connections

18

Accepting Connections: accept

```
int accept(int sockfd, struct sockaddr *cliaddr,  
socklen_t *addrlen);
```

- Accepts a client connection
 - Called by a TCP server after listen
- Parameters:
 - sockfd: listening (server) socket file descriptor
 - cliaddr: Client's address structure
 - addrlen: Length of address structure

19

Accepting Connections: accept

- Returns a new socket file descriptor
 - Corresponds to the TCP connection with the client
 - All communication with client happens on this new connection
- Listening socket is used only for accepting new connections
- If no new connection, server blocks on accept

20

Connecting to a Server: connect

```
int connect(int sockfd, struct sockaddr
            *servaddr, socklen_t *addrlen);
```

- Called by a TCP client
- Connects to a remote server at specified address and port
- Parameters:
 - sockfd: socket file descriptor
 - servaddr: Server's address structure
 - addrlen: Length of address structure

21

TCP Socket I/O

- read, write
- Remember that the return value may be different than num_bytes specified
- Should not use stream operations such as fprintf, fread, etc.
 - Buffering may cause problems

22

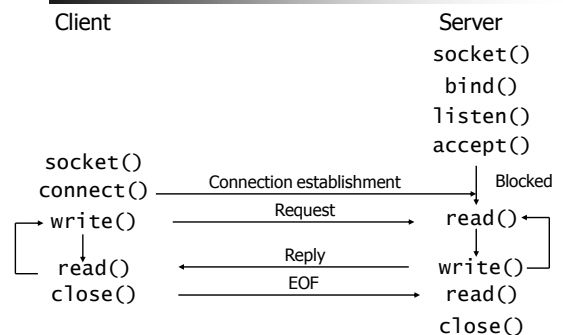
Closing a Connection: close

```
int close(int sockfd);
```

- Similar to file close
- In addition:
 - Closes TCP connection
 - Sends out any pending data before closing

23

TCP Client-Server Operations



24

Handling Server Concurrency

- TCP Server has to do multiple things
 - Listen for new connections
 - Service existing client requests
 - Perform I/O on existing client connections
- Approach 1: Iterative Server
 - Do one thing at a time
 - Accept a connection, service request, close connection, go back to waiting for new connections

25

Concurrent TCP Server

- Use processes/threads for concurrency
- Main process/thread
 - Wait for new connections
 - Accept a new connection and pass on to a worker process/thread
 - Go back to waiting
- Worker processes/threads
 - Receive client connection from main process/thread
 - Service client request, perform I/O
 - Close client connection
- Can also use asynchronous I/O for concurrency

26

UDP Sockets

- UDP is connectionless
- No connection established between client-server
- Data is transmitted as datagrams instead of stream of bytes

27

UDP Sockets: Server Operations

- Create a socket
- Bind a local address/port number
- Wait for requests from clients
- Read request, service it, return results
- No need for listen, accept

28

UDP Sockets: Client Operations

- Create a socket
- Send request to a remote server
- Receive results
- Close socket
- No need for connect, but we can call connect:
 - Associates a remote address/port pair with the socket
 - All data sent to this recorded address
 - No connection established (unlike TCP)

29

UDP Socket I/O

- Send, receive data to/from remote host

```
ssize_t sendto(int sockfd, void *buf, size_t nbytes,  
              int flags, struct sockaddr *to, socklen_t addrlen);
```

```
ssize_t recvfrom(int sockfd, void *buf, size_t nbytes,  
                int flags, struct sockaddr *from, socklen_t *addrlen);
```

- Parameters:
 - Similar to read/write
 - flags: Describe msg options (0 by default)
 - Remote address and address length

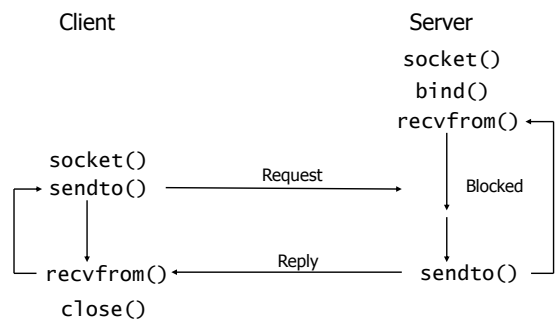
30

UDP Socket I/O

- sendto:
 - Similar to connect+write
 - No connection established
- recvfrom:
 - Similar to accept+read
 - Process blocks until a datagram is received
 - No connection established
 - Remote address structure can be NULL

31

UDP Client-Server Operation



32

Sockets Summary

- Socket Overview
- Socket Types:
 - TCP vs. UDP
 - Client vs. Server
- Socket Operations:
 - socket, bind, connect, listen, accept, ...
- Client-Server Operations
 - TCP/UDP examples
 - Concurrent TCP server